

The Legion Programming Model

Alex Aiken

Elliott Slaughter

Stanford University & SLAC

What Do You Need Today?

- A laptop
 - With access to the conference wifi
- A shell & ssh
- Login credentials
 - You should already have received these
 - But we can also give you credentials now

Acknowledgments

- The Legion project is joint work between Stanford, CMU, Los Alamos National Lab, SLAC, and NVIDIA.
- Funding has come from many sources, including NSF and DARPA, but particularly the DOE and the leadership class facilities.

Overview

Legion & Regent

- *Legion* is
 - a C++ runtime
 - a programming model
- *Regent* is a programming language
 - For the Legion programming model
 - Current implementation is embedded in Lua
 - Has an optimizing compiler
- This tutorial focuses on Regent

Regent/Legion Design Goals

- Sequential semantics
 - The better to understand what you write
 - Parallelism is extracted automatically
- Throughput-oriented
 - The latency of a single thread/process is (mostly) irrelevant
 - The overall time is what matters
- Runtime decision making
 - Because machines are unpredictable/dynamic
- Compositional
 - To enable writing parallel/distributed libraries

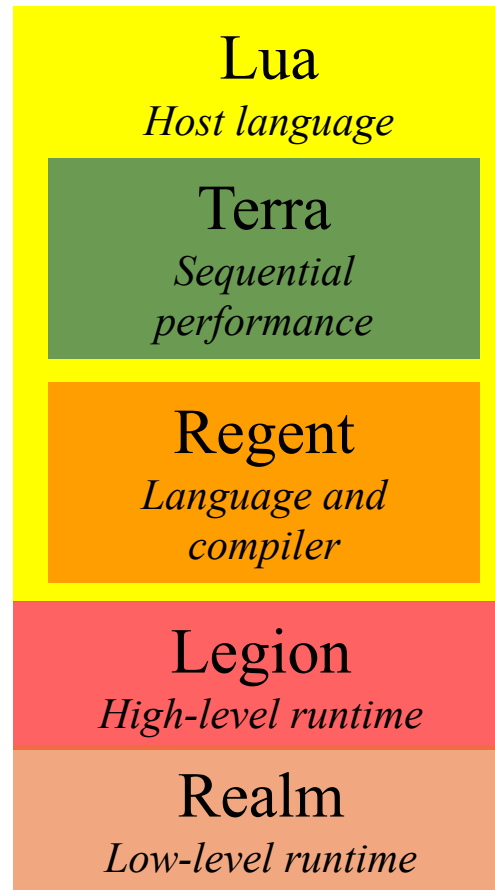
Throughput-Oriented

- Keep the machine busy
- How? Ideally,
 - Every core has a queue of independent work to do
 - Every memory unit has a queue of transfers to do
 - At all times

Consequences

- Highly asynchronous
 - Minimize synchronization
 - Esp. global synchronization
- Sequential semantics but support for parallelism
- Emphasis on describing the structure of data
 - Later

Regent Stack



Regent in Lua

- Embedded in Lua
 - Popular scripting language in the graphics community
- Excellent interoperation with C
 - And with other languages
- Python-ish syntax
 - For both Lua and Regent

- Examples Overview/1.rg & 2.rg
- To run:
 - ssh USER@sapling.stanford.edu
 - cd sc24
 - source env.sh
 - cd Overview
 - sbatch r1.sh
 - queue --me # wait until job completes
 - less slurm-*.log

Tasks

Tasks

- Tasks are Regent's unit of parallel execution
 - Distinguished functions that can be executed asynchronously
- No preemption
 - Tasks run until they block or terminate
 - And ideally they don't block ...

Blocking

- *Blocking* means a task cannot continue
 - So the task stops running
- Blocking does not prevent independent work from being done
 - If the processor has something else to do
 - Does prevent the task from continuing and launching more tasks
- Avoid blocking

Subtasks

- Tasks can call subtasks
 - Nested parallelism
- Terminology: *parent* and *child* tasks

Example

```
task tester(sum: int64)
```

```
...
```

```
end
```

```
task main()
```

```
    var sum: int64 = summer(10)
```

```
    sum = tester(sum)
```

```
    format.println("The answer is: {}",sum)
```

```
end
```


If a parent task inspects the result of a child task, the parent task blocks pending completion of the child task.

- Examples Tasks/1.rg & 2.rg

- Reminder:

```
cd sc24/Tasks
```

```
sbatch r1.sh
```

```
queue --me
```

Legion Prof

Legion Prof

- A tool for showing performance timeline
 - Each processor is a timeline
 - Each operation is a time interval
 - Different kinds of operations have different colors
- White space = idle time

Example 1: Legion Prof

```
cd sc24/Tasks
```

```
sbatch rp1.sh
```

```
squeue --me # wait for job to complete
```

```
legion_prof_to_public_html prof1_*.gz
```

Open the URL printed to the terminal

Example 2: Legion Prof

```
cd sc24/Tasks  
sbatch rp2.sh  
queue --me # wait for job to complete  
legion_prof_to_public_html prof2_*.gz
```

Open the URL printed to the terminal

Mapping

- How does Regent/Legion decide on which processor to run tasks?
- This decision is under the *mapper*'s control
- Here we are using the default mapper
 - Passes out tasks to CPUs that are not busy
 - Programmers can write their own mappers
 - More on mapping later

Parallelism

Example Tasks/3.rg

- “for all” style parallelism
- Note the order of completion of the tasks
 - `main()` finishes first (or almost first)!
 - All subtasks managed by the runtime system
 - Subtasks execute in non-deterministic order
- How?
 - Regent notices that the tasks are *independent*
 - No task depends on another task for its inputs

Runtime Dependence Analysis

- Example Tasks/4.rg is more involved
 - Positive tasks (print a positive integer)
 - Negative tasks (print a negative integer)
- Some tasks are dependent
 - The task for -5 depends on the task for 5
 - Note loop in `main()` does *not* block on the value of `j`!
- Some are independent
 - Positive tasks are independent of each other
 - Negative tasks are independent of each other


Workflow

- Use Legion Prof to find idle time
 - white space
- The profile can also show why a task ran when it did
 - What was the task waiting for?!

Exercise 1

Computing the Area of a Unit Circle

- A Monte Carlo simulation to compute the area of a unit circle inscribed in a square
- Throw darts
 - Fraction of darts landing in the circle = ratio of circle's area to square's area



Computing the Area of a Unit Circle

- Example Pi/1.rg
 - Slow!
 - Why?

Exercise 1

- Modify Pi/1.rg
 - Edit x1.rg
 - Make multiple trials per subtask
- Use
 - 4 subtasks
 - 2500 trials per subtask
- Produce profiling output
 - Run `legion_prof_to_public_html` on the resulting profile logs
- Lesson: Task *granularity* is important
 - Tasks that are too short cost more for the runtime system to analyze than to execute!
 - For Legion, an average task granularity of 1ms is recommended

Regions

Regions

- A region is a (typed) collection
- Regions are the cross product of
 - An *index space*
 - A *field space*

Regions/1.rg

Bit	
0	false
1	false
2	false
3	false
4	false
5	true
6	true
7	true
8	true
9	true

Discussion

- Regions are *the* way to organize large data collections in Regent
- Index spaces can be dense or sparse
- Any number of fields
- Built-in support for 1D, 2D and 3D index spaces

Privileges

- A task that takes region arguments must
 - Declare its *privileges* on the region
 - Reads, Writes, Reduces (with a specific reduction operator)
- The task may only perform operations for which it has privileges
 - Including any subtasks it calls

- Example Regions/2.rg
- Example Regions/3.rg

Reduction Privileges

- Regions/4.rg
 - A sequence of tasks that increment elements of a region
 - With Read/Write privileges
- Regions/5.rg
 - 4.rg but with Reduction privileges
- Note: Reductions can create additional copies
 - To get more parallelism
 - Under mapper control
 - Not always preferred to Read/Write privileges

Partitioning

Partitioning

- To enable parallelism on a region, *partition* it into smaller pieces
 - And then run a task on each piece
- Regent has a rich set of partitioning primitives

Partitioning Example

Bit

0	false
1	false
2	false
3	false
4	false
5	true
6	true
7	true
8	true
9	true

Partitioning Example

		Bit
bit_region_partition[0]	0	false
	1	false
	2	false
	3	false
	4	false
bit_region_partition[1]	5	true
	6	true
	7	true
	8	true
	9	true


Equal Partitions

- One commonly used primitive is to split a region into a number of (nearly) equal size subregions
- Partitioning/1.rg
- Partitioning/2.rg

Discussion

- Partitioning does not create copies
 - It names subsets of the data
- Partitioning does not remove the parent region
 - It still exists and can be used
- Regions and partitions are first-class values
 - Can be created, destroyed, stored in data structures, passed to and returned from tasks

Region Trees



More Discussion

- The same data can be partitioned multiple ways
 - Again, these are just names for subsets
- Subregions can themselves be partitioned

Dependence Analysis

- Regent uses tasks' region arguments to compute which tasks can run in parallel
 - What region is being accessed
 - Does it overlap with another region that is in use?
 - What field is being accessed
 - If a task is using an overlapping region, is it using the same field?
 - What are the privileges?
 - If two tasks are accessing the same field, are they both reading or both reducing?

A Crucial Fact

- Regent analyzes *sibling* tasks
 - Tasks launched directly by the same parent task
- Theorem: Analyzing dependencies between sibling tasks is sufficient to guarantee sequential semantics
- Never check for dependencies otherwise
 - Crucial to the overall design of Regent

Consequences

- Dependence analysis is a source of runtime overhead
- Can be reduced by reducing the number of sibling tasks
 - Group some tasks into subtasks
- But beware!
 - This may also reduce the available parallelism
- Partitioning/3.rg

Partitioning/3.rg

- Note that passing a region to a task does not mean the data is copied to where that task runs
 - C.f., [launcher](#) task must name the parent region for type checking reasons
- If the task doesn't touch a region/field, that data doesn't need to move


Fills

- A better way to initialize regions is to use *fill* operations

`fill(region.field, value)`

- Partitioning/4.rg

Multiple Partitions



Discussion

- Different views onto the same data
- Again, can have multiple views in use at the same time
- Regent will figure out the data dependencies

Exercise 2

- Modify Partitioning/x2.rg to
- Have two partitions of bit_region
 - One with 3 subregions of size 20
 - One with 6 subregions of size 10
- In a loop, alternately launch subtasks on one partition and then the other
- Edit x2.rg

Aliased Partitions

- So far all of our examples have been *disjoint partitions*
- It is also possible for partitions to be *aliased*
 - The subregions overlap
- Partitioning/5.rg

Partitioning Summary

- Significant Regent applications have interesting region trees
 - Multiple views
 - Aliased partitions
 - Multiple levels of nesting
- And complex task dependencies
 - Subregions, fields, privileges
 - And *coherence*, which we have not discussed
- Regions express locality
 - Data that will be used together
 - An example of a “local address space” design
 - Tasks can only access their region arguments

Image Blur

Index Notation

- First example with a 2D region
- Rect2d type
 - 2D rectangle
 - To construct: `rect2d { lo, hi }`
 - Note `lo` and `hi` are 2D points!
 - Fields: `r.lo`, `r.hi`
 - Operation: `r.lo + {1,1}`, `r.hi - {1,1}`
- The following works (modulo bounds):

```
for x in r do
  r[x + {1,1}]
```

Blur

- Compute a Gaussian blur of an image
- Edit Blur/blur.rg
 - Search for TODO
 - ... in two separate places ...
 - Test with blur.sh
- Solution is in blur_solution.rg
 - Also scripts for running the solution
 - With and without GPUs


Regions Review

- A region is a (typed) collection
- Regions are the cross product of
 - An *index space*
 - A *field space*

Region References

- Elements of a field can be indices into another region
 - The moral equivalent of a region pointer or reference
- Regent's type system tracks fields that point to other regions
 - E.g., `Edge(Nodes)` is the type of edges (indices) into a region of Nodes

RegionReferences/1.rg and 2.rg



Partitioning By Field

- A field can be used to create a partition
 - Elements with the same field value are assigned to the same subregion
 - Similar to a “group by” operation in databases
- Write elements of an index space `colors` into the field `f`
 - Using an arbitrary computation
- Then call `partition(region.f, colors)`
 - `RegionReferences/3.rg`

Dependent Partitioning

Partitioning, Revisited

- Why do we want to partition data?
 - For parallelism
 - We will launch many tasks over many subregions
- A problem
 - We often need to partition multiple data structures in a consistent way
 - E.g., given that we have partitioned the nodes a particular way, that will dictate the desired partitioning of the edges

Dependent Partitioning

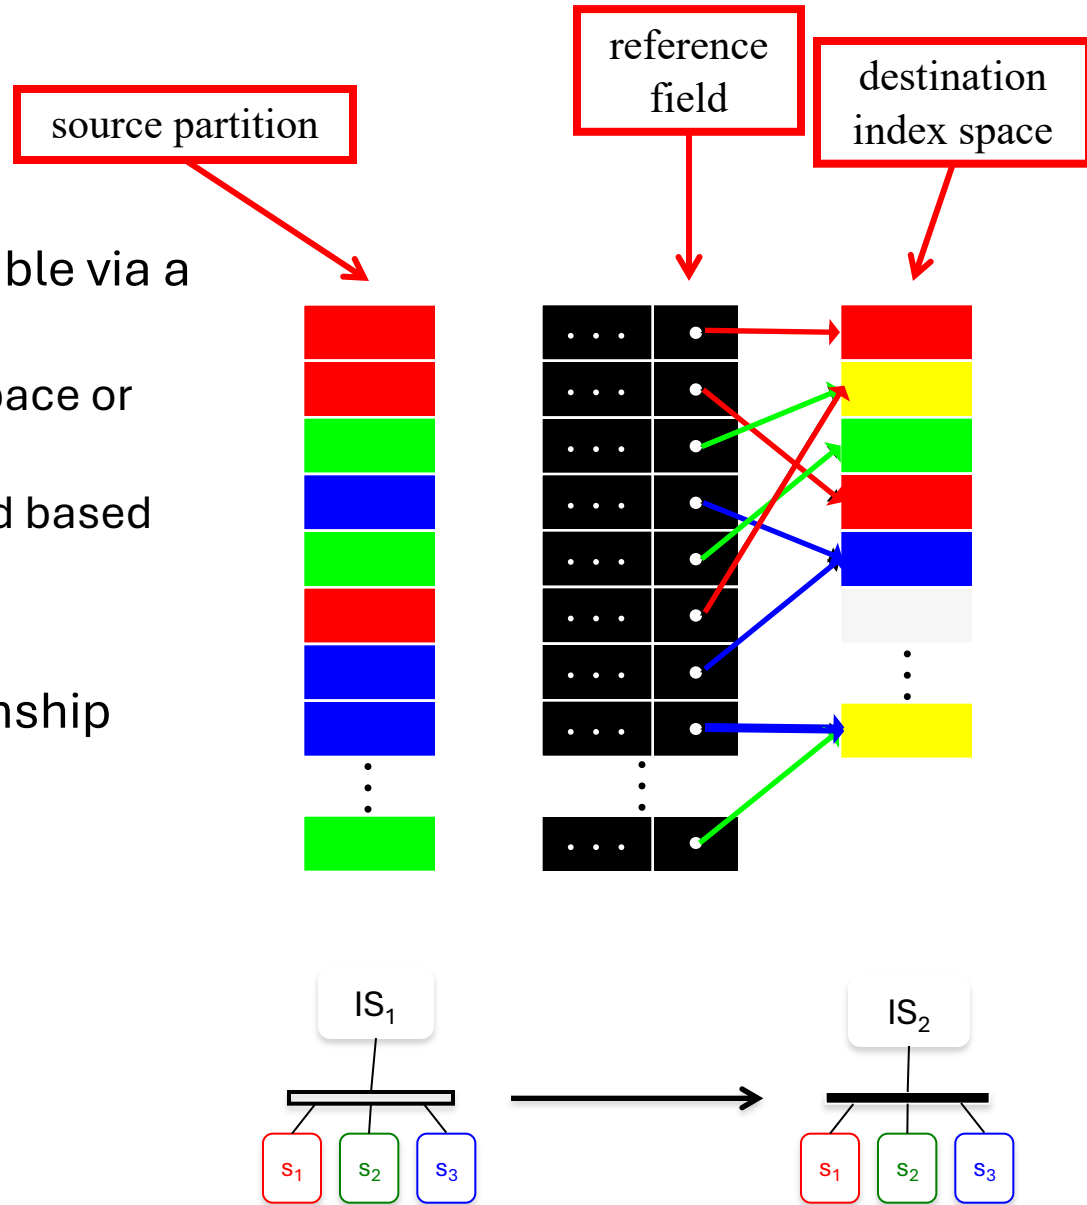
- Distinguish two kinds of partitions
- *Independent partitions*
 - Computed from the parent region, using, e.g.,
 - `partition(equal, ...)`
- *Dependent partitions*
 - Computed using another partition

Dependent Partitioning Operations

- Image
 - Use the image of a field in a partition to define a new partition
- Preimage
 - Use the pre-image of a field in a partition ...
- Set operations
 - Form new partitions using the intersection, union, and set difference of other partitions

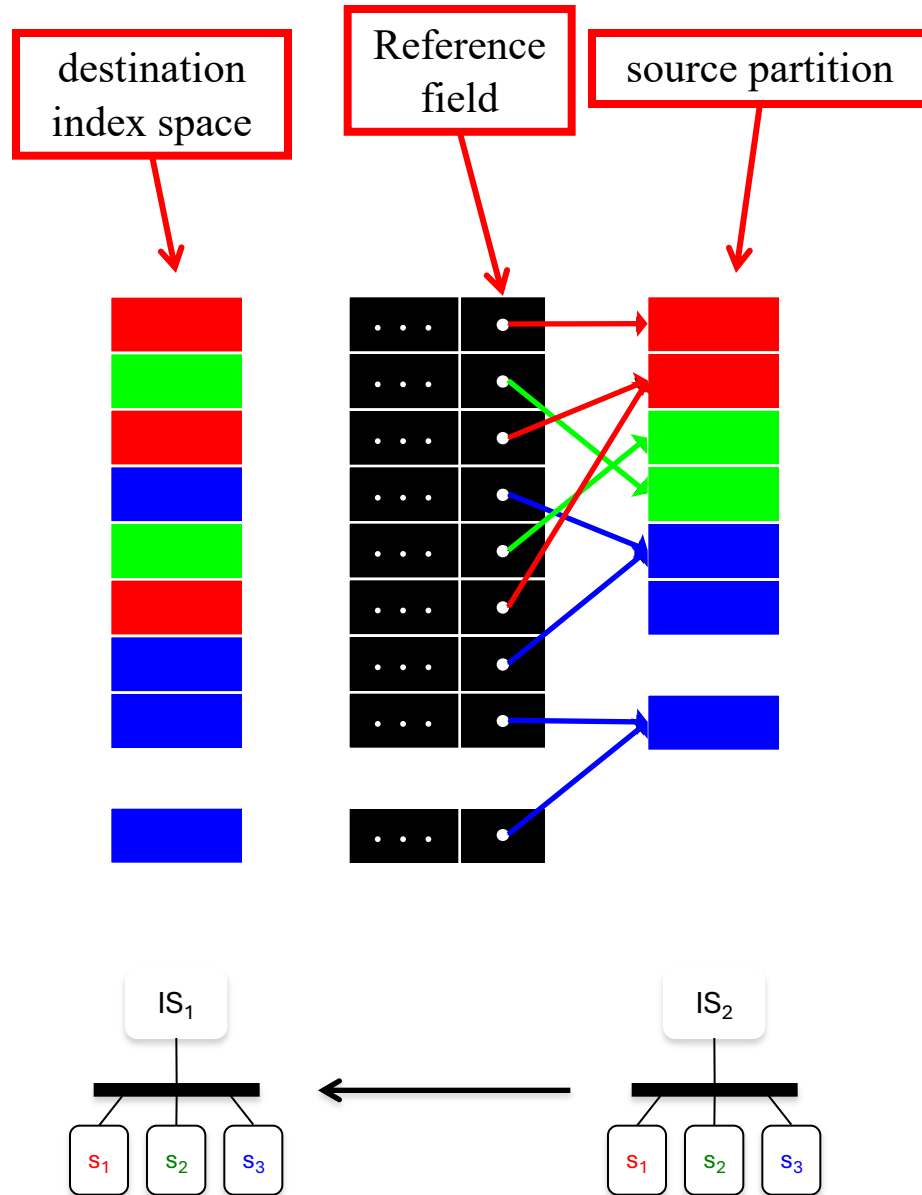
Image

- Computes elements reachable via a field lookup
 - Can be applied to index space or another partition
 - Computation is distributed based on location of data
- Regent understands relationship between partitions



Preimage

- Inverse of image
 - Computes elements that reach a given subspace
 - Preserves disjointness
- Multiple images/preimages can be combined
 - Can capture complex task access patterns



DependentPartitioning/1.rg

- Partition the nodes
 - Equal partitioning
- Then partition the edges
 - Preimage of the source node of each edge
- For each node subregion r , form a subregion of those edges where the source node is in r

DependentPartitioning/2.rg

- Partition the edges
 - Equal partitioning
- Then partition the nodes
 - Image of the source node of each edge
- For each edge subregion r , form a subregion of those nodes that are source nodes in r

Discussion

- Note that these two examples compute (almost) the same partition
- Can derive the node partition from the edges, or vice versa

Exercise

- What would the example look like if we partitioned based on the destination node?
- Let's find out ...
 - Modify 2.rg to partition using the destination node
 - Code is in DependentPartitioning/x3.rg

Set Operations: Set Difference

- Partition the edges
 - Equal partition
- Compute the source and destination node partitions of the previous two examples
- The final node partition is the set difference
 - What does this compute?
 - Examples DependentPartitioning/4.rg & 5.rg

Set Operations: Set Intersection

- Partition the edges
 - Equal partition
- Compute the source & destination node partitions
- Final node partition is the intersection
 - What does this compute?
 - Example `DependentPartitioning/6.rg`

DependentPartitioning/7.rg

- Same as the last example
- Once the final node partition is computed, compute a partition of the edges such that each edge subregion has only the edges connecting the nodes in the corresponding node subregion

Page Rank

The Algorithm

- The page rank algorithm computes an iterative solution to the following equation, where
 - $PR(p)$ is the probability that page p is visited
 - N is the number of pages
 - $L(p)$ is the number of outgoing links from p
 - d is a “damping factor” between 0 and 1

$$PR(p) = \frac{1 - d}{N} + d \sum_{p' \in M(p)} \frac{PR(p')}{L(p')}$$

Exercise

- Modify Pagerank/pagerank.rg
- Play with the partitioning of the graph
 - Can you switch from a page-based partitioning to a link-based partitioning?
- And possibly the permissions
 - See “TODO”

Control Replication

Subtasks and Scalability

Consider a task

task t(...)

...

for color in colors do

 subtask(P[color])

end

...

end

If $|\text{colors}| > \sim 32$, then simply launching the subtasks will be a sequential bottleneck.

Why? The subtasks can run in parallel, but **t** is sequential.

As $|\text{colors}|$ increases, the time to launch the subtasks becomes significant.

Subtasks and Scalability

Consider a task

```
task t(...)
```

```
...
```

```
  for color in colors do
```

```
    subtask(P[color])
```

```
  end
```

```
...
```

```
end
```

Idea: Spread **t**'s work over multiple processors.

Each launches a subset of the subtasks.

An Issue ...

Consider a task

```
task t(...)
  ...
  i += 1
  var r = region(ispace, fspace)
  format.println("Hello!")

  for color in colors do
    subtask(P[color])
  end
end
```

Must preserve **t**'s semantics!

If **t** is executed cooperatively by multiple threads, the effect must still be as if **t** were executed sequentially.

- Index space/region/partition creation happens once
- Print statements happen once

But where possible code should be replicated SPMD-style to avoid communication and synchronization

- Every thread can increment its own local copy of **i**

Control Replication

To execute a task `t` cooperatively on multiple processors:

```
__demand(__replicable)  
task t(...) ... end
```

The runtime will replicate the task on multiple threads



- With each thread handling a subset of the subtask launches
- And replicating the rest of the task, where it is safe to do so
- `ControlReplication/1.rg`

Mapping

Mapping

- Mapping is the process of assigning resources to Regent/Legion programs
- The runtime issues and maps tasks *ahead of actual execution*
 - In a well-structured program, mapping decisions are made well before it is time to run a task

Picture



Mapping

- Mapping is the process of assigning resources to Regent/Legion programs
- Most importantly
 - Assign a processor to each task
 - The task will execute in its entirety on that processor
 - Assign a memory to each region argument

Mappers

- Legion provides a mapping API
 - An interface for making mapping decisions
- The mapping API requires a standard set of functions
 - An implementation of these functions is a *mapper*
- Legion comes with a default mapper
 - A set of heuristics for mapping that makes reasonable decisions
 - But for highest performance applications generally require custom mappers

Task Variants

- A task can have multiple *variants*
 - Different implementations of the same task
 - Multiple variants can be registered with the Legion runtime
- Examples
 - A variant for CPUs
 - Another variant for GPUs
 - Variants for different data layouts

Controlling Processor Mapping in Regent

- By default, the Regent compiler produces only CPU task variants
 - So the default mapper runs these tasks on CPUs
- Place immediately before a task declaration
 - `__demand(__cuda)`
 - Causes both CPU and GPU task variants to be produced
 - The default mapper prefers a GPU variant if one exists

Where is the Data Placed?

- The default mapper ...
- ... places region arguments in CPU RAM if the corresponding task is mapped to a CPU
- ... places region arguments in the frame buffer if the corresponding task is mapped to a GPU
- There are other possible memories (e.g., zero copy memory) that can be used with a custom mapper

Scratching the Surface

- There are many other mapping options
- Data layout in regions
- Eager/lazy collection of dead regions
- Backpressure – how far ahead of execution should the runtime get?
- ...

Libraries

Libraries

- One of the design goals of Legion/Regent is to enable the development of distributed/parallel libraries
 - Composable for both functionality and performance
- Task-based systems are natural for library development
 - Composing independently written tasking programs always makes sense

cuNumeric

- A library from NVIDIA
- Maps NumPy programs to Legion
 - NumPy arrays become regions
 - NumPy array operations are partitioned across the machine

URL: <https://developer.nvidia.com/cunumeric>

Checkpointing

- A library from SLAC
- Provides checkpointing
 - Low overhead
 - Easy to use
 - in Regent, just add `__checkpoint`
 - in C++, requires a header and call `runtime->checkpoint(...)`

<https://github.com/StanfordLegion/resilience/>

FlexFlow

- A deep learning framework
- Automatically parallelizes deep learning training
 - Searches for a good partitioning of the data
 - Handles model, tensor, data parallelism
- URL: <https://flexflow.ai>

Research Libraries

- Dense tensor algebra
 - Sparse tensor algebra
 - Krylov solvers
 - FFT
 - Mapping DSL
-
- Happy to work with anyone interested in using these!

Conclusions

Conclusions

- Legion/Regent is a task-based parallel programming system
- Advantages
 - Easy to exploit multiple levels of parallelism in a uniform manner
 - Separate machine mapping
 - Novel and rich partitioning sublanguage
 - Designed to support writing (distributed) libraries
- Good/great performance and portability!